



EECS 230 Deep Learning

Lecture 2: Machine Learning


Some slides from Olga Veksler and Simon Prince



Linear Algebra

What is linear algebra?

- Linear algebra is the branch of mathematics concerning linear equations such as
$$a_1x_1 + \dots + a_nx_n = b$$
 - In vector notation we say $a^T x = b$
 - Called a linear transformation of x
- Linear algebra is fundamental to geometry, for defining objects such as lines, planes, rotations



Linear equation $a_1x_1 + \dots + a_nx_n = b$ defines a plane in (x_1, \dots, x_n) space
Straight lines define common solutions to equations

Linear Algebra Topics

- ❑ Scalars, Vectors, Matrices and Tensors
- ❑ Multiplying Matrices and Vectors
- ❑ Identity and Inverse Matrices
- ❑ Linear Dependence and Span
- ❑ Norms
- ❑ Special kinds of matrices and vectors
- ❑ Eigendecomposition
- ❑ Singular value decomposition
- ❑ The Moore Penrose pseudoinverse
- ❑ The trace operator
- ❑ The determinant
- ❑ Ex: principal components analysis



Scalar

- Single number
 - In contrast to other objects in linear algebra, which are usually arrays of numbers
- Represented in lower-case italic x
 - They can be real-valued or be integers
 - E.g., let $x \in \mathbb{R}$ be the slope of the line
 - Defining a real-valued scalar
 - E.g., let $n \in \mathbb{N}$ be the number of units
 - Defining a natural number scalar

Vector

- An array of numbers arranged in order
- Each no. identified by an index
- Written in lower-case bold such as \mathbf{x}
 - its elements are in italics lower case, subscripted

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix}$$

- If each element is in R then \mathbf{x} is in R
- We can think of vectors as points in space

Matrices

- 2-D array of numbers
 - So each element identified by two indices
- Denoted by bold typeface A
 - Elements indicated by name in italic but not bold
 - $A_{1,1}$ is the top left entry and $A_{m,n}$ is the bottom right entry
 - We can identify nos in vertical column j by writing : for the horizontal coordinate
- E.g.,
$$A = \begin{bmatrix} A_{1,1} & A_{1,2} \\ A_{2,1} & A_{2,2} \end{bmatrix}$$
 - $A_{i:}$ is i^{th} row of A , $A_{:j}$ is j^{th} column of A
- If A has shape of height m and width n with real-values then $A \in R^{m \times n}$

Tensor

- Sometimes need an array with more than two axes
 - E.g., an RGB color image has three axes
- A tensor is an array of numbers arranged on a regular grid with variable number of axes
 - See figure next
- Denote a tensor with this bold typeface: \mathbf{A}
- Element (i,j,k) of tensor denoted by $A_{i,j,k}$

Multiplying matrices

- For product $C = AB$ to be defined, A has to have the same no. of columns as the no. of rows of B
 - If A is of shape $m \times n$ and B is of shape $n \times p$ then *matrix product* C is of shape $m \times p$

$$C = AB \Rightarrow C_{i,j} = \sum_k A_{ik} B_{kj}$$

- Note that the standard product of two matrices is not just the product of two individual elements
 - Such a product does exist and is called the element-wise product or the Hadamard product $A \odot B$

Linear transformation

- $A\mathbf{x} = \mathbf{b}$
 - where $A \in \mathbb{R}^{n \times n}$ and $\mathbf{b} \in \mathbb{R}^n$

– More explicitly

$$A_{11}x_1 + A_{12}x_2 + \dots + A_{1n}x_n = b_1$$

$$A_{21}x_1 + A_{22}x_2 + \dots + A_{2n}x_n = b_2$$

$$A_{n1}x_1 + A_{n2}x_2 + \dots + A_{nn}x_n = b_n$$

n equations in
 n unknowns

$$A = \begin{bmatrix} A_{1,1} & \cdots & A_{1,n} \\ A_{n,1} & \ddots & A_{nn} \end{bmatrix} \quad \mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} \quad \mathbf{b} = \begin{bmatrix} b_1 \\ b_2 \\ \vdots \\ b_n \end{bmatrix}$$

$n \times n$ $n \times 1$ $n \times 1$

Can view A as a *linear transformation* of vector \mathbf{x} to vector \mathbf{b}

- Sometimes we wish to solve for the unknowns $\mathbf{x} = \{x_1, \dots, x_n\}$ when A and \mathbf{b} provide constraints

Matrix inverse

- Inverse of square matrix A defined as $A^{-1}A=I_n$
- We can now solve $Ax=b$ as follows:

$$Ax=b$$

$$A^{-1}Ax=A^{-1}b$$

$$I_n x = A^{-1} b$$

$$x = A^{-1} b$$

- This depends on being able to find A^{-1}
- If A^{-1} exists there are several methods for finding it

Norms

- Used for measuring the size of a vector
- Norms map vectors to non-negative values
- Norm of vector $\mathbf{x}=[x_1, \dots, x_n]^T$ is distance from origin to \mathbf{x}
 - It is any function f that satisfies:

$$f(\mathbf{x}) = 0 \Rightarrow \mathbf{x} = 0$$

$$f(\mathbf{x} + \mathbf{y}) \leq f(\mathbf{x}) + f(\mathbf{y}) \quad \text{Triangle Inequality}$$

$$\forall \alpha \in R \quad f(\alpha \mathbf{x}) = |\alpha| f(\mathbf{x})$$



L^p Norm

- Definition:

- L^2 Norm

- Called Euclidean norm
 - Simply the Euclidean distance between the origin and the point \mathbf{x}
 - written simply as $\|\mathbf{x}\|$
 - Squared Euclidean norm is same as $\mathbf{x}^T \mathbf{x}$



- L^1 Norm

- Useful when 0 and non-zero have to be distinguished
 - Note that L^2 increases slowly near origin, e.g., $0.1^2=0.01$

- L^∞ Norm

- Called max norm

$$\|\mathbf{x}\|_p = \left(\sum_i |x_i|^p \right)^{\frac{1}{p}}$$



$$\|\mathbf{x}\|_\infty = \max_i |x_i|$$

Special kind of vectors

- Unit Vector
 - A vector with unit norm


$$\|x\|_2 = 1$$

$$\begin{bmatrix} 2 \\ -3 \\ -2 \end{bmatrix}, \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 4 \\ -2 \\ 7 \end{bmatrix}$$

- Orthogonal Vectors

- A vector x and a vector y are orthogonal to each other if $x^T y = 0$

- If vectors have nonzero norm, vectors at 90 degrees to each other



- Orthonormal Vectors

- Vectors are orthogonal & have unit norm
 - Orthogonal Matrix

- A square matrix whose rows are mutually

- orthonormal: $A^T A = A A^T = I$


$$A^{-1} = A^T$$

Eigenvector

- An eigenvector of a square matrix A is a non-zero vector v such that multiplication by A only changes the scale of v

$$Av = \lambda v$$

- The scalar λ is known as eigenvalue
- If v is an eigenvector of A , so is any rescaled vector sv . Moreover sv still has the same eigen value. Thus look for a unit eigenvector



Wikipedia

Eigendecomposition


- Suppose that matrix A has n linearly independent eigenvectors $\{\mathbf{v}^{(1)}, \dots, \mathbf{v}^{(n)}\}$ with eigenvalues $\{\lambda_1, \dots, \lambda_n\}$
- Concatenate eigenvectors to form matrix V
- Concatenate eigenvalues to form vector $\lambda = [\lambda_1, \dots, \lambda_n]$
- Eigendecomposition of A is given by

$$A = V \text{diag}(\lambda) V^{-1}$$

Effect of eigenvalue and eigenvector


- Example of 2×2 matrix
- Matrix A with two orthonormal eigenvectors
 - $v^{(1)}$ with eigenvalue λ_1 , $v^{(2)}$ with eigenvalue λ_2

Plot of unit vectors $u \in U^2$
(circle)



with two variables x_1 and x_2

Plot of vectors Au
(ellipse)



Positive Semidefinite Matrix (PSD)

- A matrix whose eigenvalues are all positive is called *positive definite*
 - Positive or zero is called *positive semidefinite*
- If eigen values are all negative it is *negative definite*
 - Positive definite matrices guarantee that $\mathbf{x}^T \mathbf{A} \mathbf{x} \geq 0$

Singular Value Decomposition (SVD)

- Eigendecomposition has form: $A=V\text{diag}(\lambda)V^{-1}$
 - If A is not square, eigendecomposition is undefined
- SVD is a decomposition of the form $A=UDV^T$
- SVD is more general than eigendecomposition
 - Used with any matrix rather than symmetric ones
 - Every real matrix has a SVD
 - Same is not true of eigen decomposition



Probability and Statistics

Probability and Statistics

- Probability Theory
 - A mathematical framework for representing uncertain statements
 - Provides a means of quantifying uncertainty and axioms for deriving new uncertain statements
- Use of probability theory in artificial intelligence
 - 1.Tells us how AI systems should reason
 - So we design algorithms to compute or approximate various expressions using probability theory
 - 2.Theoretically analyze behavior of AI systems

Random Variable

- Variable that can take different values randomly
- Scalar random variable denoted x
- Vector random variable is denoted in bold as \mathbf{x}
- Values of r.v.s denoted in italics x or \mathbf{x}
 - Values denoted as $\text{Val}(x)=\{x_1, x_2\}$
- Random variable must has a probability distribution to specify how likely the states are
- Random variables can be discrete or continuous
 - Discrete values need not be integers, can be named states
 - Continuous random variable is associated with a real value

Probability Distribution

- ❑ A probability distribution is a description of how likely a random variable or a set of random variables is to take each of its possible states
- ❑ The way to describe the distribution depends on whether it is discrete or continuous

Continuous Variables and PDFs

- When working with continuous variables, we describe probability distributions using probability density functions
- To be a pdf p must satisfy:
 - The domain of p must be the set of all possible states of x .
 - $\forall x \in \mathbf{x}, p(x) \geq 0$. Note that we do not require $p(x) \leq 1$.
 - $\int p(x)dx = 1$.

Marginal distribution

- ❑ Sometimes we know the joint distribution of several variables
- ❑ And we want to know the distribution over some of them
- ❑ It can be computed using

$$\forall x \in \mathbf{x}, P(\mathbf{x} = x) = \sum_y P(\mathbf{x} = x, \mathbf{y} = y)$$

$$p(x) = \int p(x, y) dy$$

Conditional probability

- We are often interested in the probability of an event given that some other event has happened
- This is called conditional probability
- It can be computed using

$$P(y = y \mid x = x) = \frac{P(y = y, x = x)}{P(x = x)}.$$

Chain rule of conditional probability

- Any probability distribution over many variables can be decomposed into conditional distributions over only one variable

$$P(x^{(1)}, \dots, x^{(n)}) = P(x^{(1)}) \prod_{i=2}^n P(x^{(i)} \mid x^{(1)}, \dots, x^{(i-1)})$$

- An example with three variables

$$P(a, b, c) = P(a \mid b, c)P(b, c)$$

$$P(b, c) = P(b \mid c)P(c)$$

$$P(a, b, c) = P(a \mid b, c)P(b \mid c)P(c)$$

Independence and conditional independence

- Independence: $x \perp\!\!\! \perp y$

– Two variables x and y are independent if their probability distribution can be expressed as a product of two factors, one involving only x and the other involving only y

$$\forall x \in X, y \in Y, p(x = x, y = y) = p(x = x)p(y = y)$$

- Conditional Independence: $x \perp\!\!\! \perp y \mid z$

– Two variables x and y are independent given variable z , if the conditional probability distribution over x and y factorizes in this way for every z

$$\forall x \in X, y \in Y, z \in Z, p(x = x, y = y \mid z = z) = p(x = x \mid z = z)p(y = y \mid z = z)$$

Common probability distribution

- Several simple probability distributions are useful in many contexts in machine learning
 - Bernoulli over a single binary random variable
 - Multinoulli distribution over a variable with k states
 - Gaussian distribution
 - Mixture distribution

Mixture of Distribution


- A mixture distribution is made up of several component distributions
- On each trial, the choice of which component distribution generates the sample is determined by sampling a component identity from a multinoulli distribution:

$$P(\mathbf{x}) = \sum_i P(c = i)P(\mathbf{x} \mid c = i)$$

– where $P(c)$ is a multinoulli distribution

Gaussian mixture model

- Components $p(x|c=i)$ are Gaussian
- Each component has a separately parameterized mean $\mu^{(i)}$ and covariance $\Sigma^{(i)}$
- Any smooth density can be approximated with enough components
- Samples from a GMM:
 - 3 components



Bayes's rule

- **Bayes' theorem** (alternatively **Bayes' law** or **Bayes' rule**), named after Thomas Bayes, describes the probability of an event, based on prior knowledge of conditions that might be related to the event.
- For example, if the risk of health problems is known to increase with age, Bayes' theorem allows the risk to an individual of a known age to be assessed more accurately by conditioning it relative to their age.

$$P(A|B) = \frac{P(A \cap B)}{P(B)} = \frac{P(A) * P(B|A)}{P(B)}$$



Machine Learning

Major Types of machine learning

- ❑ Supervised learning: Given pairs of input-output, learn to map the input to output
 - ❑ Image classification
 - ❑ Speech recognition
 - ❑ Regression (continuous output)
- ❑ Unsupervised learning: Given unlabeled data, uncover the underlying structure or distribution of the data
 - ❑ Clustering
 - ❑ Dimensionality reduction
- ❑ Reinforcement learning: training an agent to make decisions within an environment to maximize a cumulative reward
 - ❑ Game playing (e.g., AlphaGo)
 - ❑ Robot control

Subtypes of supervised ML

❑ Classification

- ❑ output belongs to a finite set
- ❑ example: $\text{age} \in \{\text{baby, child, adult, elder}\}$
- ❑ output is also called class or label


❑ Regression

- ❑ output is continuous
- ❑ examples: $\text{age} \in [0, 130]$



❑ Difference mostly in design of loss functions

Example: supervised digit recognition


- ❑ Easy to collect images of digits with their correct labels



- ❑ ML algorithm can use collected data to produce a program for recognizing previously unseen images of digits



Example: Regression



Supervised ML

- We are given

1. Training examples $\mathbf{x}^1, \mathbf{x}^2, \dots, \mathbf{x}^n$
 2. Target output for each sample y^1, y^2, \dots, y^n
- $\left. \begin{array}{l} \\ \end{array} \right\} labeled\ data$

- **Training phase**

- estimate function $y = f(\mathbf{x})$ from labeled data
 - where $f(\mathbf{x})$ is called *classifier, learning machine, prediction function, etc.*

- **Testing phase** (deployment)

- predict output $f(\mathbf{x})$ for a new (unseen) sample \mathbf{x}

Training/Testing Phases Illustrated

Training

training examples

0	0	0	0	0	0	0	0	0
1	1	1	1	1	1	1	1	1
2	2	2	2	2	2	2	2	2
3	3	3	3	3	3	3	3	3
4	4	4	4	4	4	4	4	4
5	5	5	5	5	5	5	5	5
6	6	6	6	6	6	6	6	6
7	7	7	7	7	7	7	7	7
8	8	8	8	8	8	8	8	8
9	9	9	9	9	9	9	9	9

feature vectors

training labels

Training

Learned model f

Testing



test Image

feature vector

Learned model f

label prediction

Training phase as parameter estimation

- Estimate prediction function $y = f(x)$ from labeled data

Typically, search for f is limited to some type/group of functions (“*hypothesis space*”) parameterized by *weights* w that must be estimated

$$f_w(x) \text{ or } f(w, x)$$

$$w = ?$$

Goal: find classifier parameters (weights) w so that $f(w, xi) = y^i$ “as much as possible” for all training examples,

Loss function

- Training dataset of I pairs of input/output examples

$$\{\mathbf{x}_i, \mathbf{y}_i\}_{i=1}^I$$

- Loss function or cost function measures how bad model is:

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \sum_i L(\mathbf{y}^i, \mathbf{f}(\mathbf{w}, \mathbf{x}^i))$$

- ϕ is also a common notation for weights


Example: 1D Linear regression

□ Model:

$$\begin{aligned}y &= f[x, \phi] \\&= \phi_0 + \phi_1 x\end{aligned}$$

□ Parameters

$$\phi = \begin{bmatrix} \phi_0 \\ \phi_1 \end{bmatrix} \quad \begin{array}{l} \text{y-offset} \\ \text{slope} \end{array}$$



Example: 1D Linear regression


□ Model:

$$y = f[x, \phi] \\ = \phi_0 + \phi_1 x$$


□ Parameters

$$\phi = \begin{bmatrix} \phi_0 \\ \phi_1 \end{bmatrix}$$

← y-offset
← slope



Example: 1D Linear regression training data




Loss function:

$$\begin{aligned} L[\phi] &= \sum_{i=1}^I (f[x_i, \phi] - y_i)^2 \\ &= \sum_{i=1}^I (\phi_0 + \phi_1 x_i - y_i)^2 \end{aligned}$$

"Least squares loss function"

Example: 1D Linear regression training data




Loss function:

$$\begin{aligned} L[\phi] &= \sum_{i=1}^I (f[x_i, \phi] - y_i)^2 \\ &= \sum_{i=1}^I (\phi_0 + \phi_1 x_i - y_i)^2 \end{aligned}$$

“Least squares loss function”

Example: 1D Linear regression training data




Loss function:

$$\begin{aligned} L[\phi] &= \sum_{i=1}^I (f[x_i, \phi] - y_i)^2 \\ &= \sum_{i=1}^I (\phi_0 + \phi_1 x_i - y_i)^2 \end{aligned}$$

“Least squares loss function”

Example: 1D Linear regression training data




Loss function:

$$\begin{aligned} L[\boldsymbol{\phi}] &= \sum_{i=1}^I (f[x_i, \boldsymbol{\phi}] - y_i)^2 \\ &= \sum_{i=1}^I (\phi_0 + \phi_1 x_i - y_i)^2 \end{aligned}$$

“Least squares loss function”

Example: 1D Linear regression loss function





Loss function:



$$\begin{aligned}L[\phi] &= \sum_{i=1}^I (f[x_i, \phi] - y_i)^2 \\&= \sum_{i=1}^I (\phi_0 + \phi_1 x_i - y_i)^2\end{aligned}$$

“Least squares loss function”



Example: 1D Linear regression loss function




Example: 1D Linear regression loss function



Example: 1D Linear regression loss function




Example: 1D Linear regression loss function




Gradient Descent

❑ Example: for a function of two variables




Stop at a **local minima** where $\nabla L = \vec{0}$


Example: 1D Linear regression training




Example: 1D Linear regression training





Example: 1D Linear regression training



Example: 1D Linear regression training



Example: 1D Linear regression training




Possible objections

- ❑ But you can fit the line model in closed form!
 - ❑ Yes – but we won't be able to do this for more complex models
- ❑ But we could exhaustively try every slope and intercept combo!
 - ❑ Yes – but we won't be able to do this when there are a million parameters

Example: Linear Classification

- ❑ For example: fish classification - *salmon* or *sea bass*?
- ❑ extract two features, *fish length* and *fish brightness*



- ❑ y^i is the output (label or target) for example \mathbf{x}^i

Linear classifier example: *perceptron*

m -dimensional
feature vector $\mathbf{x}^i \in \mathcal{R}^m$

with m components


$$\mathbf{x}^i = \begin{bmatrix} \mathbf{x}_1^i \\ \mathbf{x}_2^i \\ \vdots \\ \mathbf{x}_m^i \end{bmatrix}$$

sub-indices are for
feature components

while

super-indices are for
data points (feature vectors)

Frank Rosenblatt, 1958
inspired by neurons




more on the next slides


NOTE: for simplicity, we omit
super-indices (or **sub-indices**)
assuming the context is "clear"

Linear classifier example: *perceptron*

For two class problem and 2-dimensional data (feature vectors)



consider some
linear transformation
from 2D space to 1D
 $\mathbf{w}_0 + \mathbf{w}_1 x_1 + \mathbf{w}_2 x_2$




points of
two classes
can be
completely
mixed

Question:



Is it possible to find a linear transformation onto 1D so that transformed 1D points can be separated (by a *threshold*)?

Linear classifier example: *perceptron*

For two class problem and 2-dimensional data (feature vectors)



“good”
linear transformation
from 2D space to 1D
 $w_0^* + w_1^*x_1 + w_2^*x_2$




Answer:



In this case, YES, because the data is linearly separable
in the original feature space. So, what is the transformation?

Linear classifier example: *perceptron*

For two class problem and 2-dimensional data (feature vectors)



“good”
linear transformation
from 2D space to 1D
 $w_0^* + w_1^*x_1 + w_2^*x_2$




(w_1, w_2) – are the coordinates of the **normal**
 w_0 – is “bias” (shifting threshold to 0)

Answer:



This $2D \rightarrow 1D$ linear transformation is a projection onto the **normal** of the separating **hyper-plane**.

Linear classifier example: *perceptron*

For two class problem and 2-dimensional data (feature vectors)



**“good”
linear transformation
from 2D space to 1D**
 $w_0^* + w_1^*x_1 + w_2^*x_2$




(w_1, w_2) – are the coordinates of the **normal**
 w_0 – is “bias” (shifting threshold to 0)

In fact, any $2D \rightarrow 1D$ linear transformation $\mathbf{w} = (w_0, w_1, w_2)$ is a **projection onto normal of some hyper-plane**. So, original question really asks if there is a hyper-plane separating data.



Linear classifier example: *perceptron*

For two class problem and 2-dimensional data (feature vectors)




“good”
linear transformation
from 2D space to 1D

$$w_0^* + w_1^* x_1 + w_2^* x_2$$



thresholding
can be formally
represented by this
prediction function


$$f(\mathbf{w}, \mathbf{x}) = u(w_0 + w_1 x_1 + w_2 x_2) \quad f(\mathbf{w}, \mathbf{x}) \in \{0, 1\}$$



unit step function $u(t) := \begin{cases} 1 & \text{if } t > 0 \\ 0 & \text{O.W.} \end{cases}$
(a.k.a. Heaviside function)

Linear classifier example: perceptron

For two class problem and 2-dimensional data (feature vectors)




- Can be generalized to feature vectors \mathbf{x} of any dimension m :
$$f(W, X) = u(W^T X) \quad \text{for } W^T = [w_0, w_1, \dots, w_m] \text{ and } X^T = [1, x_1, x_2, \dots, x_m]$$

“bias”
homogeneous representation
of feature vector \mathbf{x}
- Classifier that makes decisions based on linear combination of features is called a **linear classifier**


Linear Classifiers

bad w



projected points onto
normal line are all mixed-up


better w



projected points onto
normal line are well separated


Underfitting

For some types of data no linear decision boundary can separate the samples well




- ❑ Classifier underfits the data if it can produce decision boundaries that are too simple for this type of data
 - chosen classifier type (hypothesis space) is not expressive enough

More complex (non-linear) classifiers




- ❑ for example, if $f(\mathbf{w}, \mathbf{x})$ is a polynomial of high degree
- ❑ can achieve **0%** classification error

More complex (non-linear) classifiers



- ❑ The goal is to classify well on **new data**
- ❑ Test “wiggly” classifier on new data: **25% error**

Overfitting




- Amount of data for training is always limited
- Complex model often has too many parameters to fit reliably to limited data
- Complex model may adapt too closely to “random noise” in training data, rather than look at a “big picture”

Overfitting: Extreme Example


- ❑ Two class problem: *face* and *non-face* images
- ❑ Memorize (i.e. store) all the “face” images
- ❑ For a new image, see if it is one of the stored faces
 - ❑ if yes, output “face” as the classification result
 - ❑ If no, output “non-face”
- ❑ **problem:**
 - ❑ zero error on stored data, 50% error on test (new) data
 - ❑ decision boundary is very irregular
- ❑ Such learning is **memorization without generalization**

Generalization

training data




new data



- Ability to produce correct outputs on previously unseen examples is called **generalization**
- Big question of learning theory: how to get good generalization with a limited number of examples
- Intuitive idea: **favor simpler classifiers**
- Simpler decision boundary may not fit ideally to training data but tends to generalize better to new data


Underfitting → Overfitting

underfitting




- high training error
- high test error

"just right"



- low training error
- low test error

overfitting



- low training error
- high test error