



# EECS 230 Deep Learning

## Lecture 15: Graph Neural Network

Some slides from Simon Prince, Paul-Edouard Sarlin, and Jure Leskovec




# Outline

- ❑ Graph Neural Network
  - ❑ Graph convolution layer
  - ❑ Graph convolutional network
  - ❑ Graph attention network
- ❑ An application to correspondence matching
  - ❑ SuperGlue for visual localization



# Graph Neural Network




# Real-world graphs



**Figure 13.1** Real-world graphs. Some objects, such as a) road networks, b) molecules, and c) electrical circuits, are naturally structured as graphs.



# Types of graphs

- a) social network is an undirected graph
- b) citation network is a directed graph
- c) Knowledge graph is a directed heterogeneous multigraph




# Types of graphs

- ❑ d) point cloud as a geometric graph
- ❑ e) Scene graph is hierarchical



# Graph representation

- ❑ A graph is defined as a tuple  $G = (V, E)$ 
  - ❑ where  $V$  is a set of nodes
  - ❑ and  $E$  is a set of edges
- ❑ An example graph with 6 nodes and 7 edges



b)

Adjacency matrix,  $\mathbf{A}$   
 $N \times N$

	1	2	3	4	5	6
1	■	■	■	■	■	■
2	■	■	■	■	■	■
3	■	■	■	■	■	■
4	■	■	■	■	■	■
5	■	■	■	■	■	■
6	■	■	■	■	■	■

c)

Node data,  $\mathbf{X}$   
 $D \times N$

1	2	3	4	5	6
■	■	■	■	■	■
■	■	■	■	■	■
■	■	■	■	■	■
■	■	■	■	■	■
■	■	■	■	■	■


d)

Edge data,  $\mathbf{E}$   
 $D_E \times E$

1	1	2	3	3	4	5
3	4	3	4	6	5	6
■	■	■	■	■	■	■
■	■	■	■	■	■	■
■	■	■	■	■	■	■
■	■	■	■	■	■	■

# Properties of adjacency matrix

◻  $A^k x$  gives the number of paths of length  $k$  to a node



b)

$$A = \begin{bmatrix} 0 & 1 & 0 & 1 & 0 & 0 & 0 & 0 \\ 1 & 0 & 1 & 1 & 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 & 1 & 0 & 0 & 0 \\ 1 & 1 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 1 & 1 & 1 & 0 & 1 & 0 & 1 \\ 0 & 0 & 0 & 0 & 1 & 0 & 1 & 1 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 1 & 0 & 0 \end{bmatrix}$$

c)


$$A^2 = \begin{bmatrix} 2 & 1 & 1 & 1 & 2 & 2 & 0 & 0 & 0 \\ 1 & 4 & 1 & 2 & 2 & 1 & 0 & 0 & 1 \\ 1 & 1 & 2 & 2 & 1 & 1 & 0 & 0 & 1 \\ 1 & 2 & 2 & 3 & 1 & 1 & 0 & 0 & 1 \\ 2 & 2 & 1 & 1 & 5 & 1 & 1 & 1 & 1 \\ 0 & 1 & 1 & 1 & 1 & 3 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 1 & 0 & 1 & 1 & 1 \\ 0 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 2 \end{bmatrix}$$

d)

$$x = \begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 1 \\ 0 \\ 0 \end{bmatrix}$$


e)

$$Ax = \begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 1 \\ 0 \\ 1 \end{bmatrix}$$




f)

$$A^2x = \begin{bmatrix} 0 \\ 1 \\ 1 \\ 1 \\ 3 \\ 0 \\ 1 \end{bmatrix}$$



# Permutation invariance

- ❑ A graph neural network should be permutation invariant
- ❑ CNN? MLP?




b) Adjacency A

	1	2	3	4	5	6
1	■	■	■	■	■	■
2	■	■	■	■	■	■
3	■	■	■	■	■	■
4	■	■	■	■	■	■
5	■	■	■	■	■	■
6	■	■	■	■	■	■

c) Node data, X

	1	2	3	4	5	6
1	■	■	■	■	■	■
2	■	■	■	■	■	■
3	■	■	■	■	■	■
4	■	■	■	■	■	■
5	■	■	■	■	■	■
6	■	■	■	■	■	■



e) Adjacency A


	1	2	3	4	5	6
1	■	■	■	■	■	■
2	■	■	■	■	■	■
3	■	■	■	■	■	■
4	■	■	■	■	■	■
5	■	■	■	■	■	■
6	■	■	■	■	■	■

f) Node data, X

	1	2	3	4	5	6
1	■	■	■	■	■	■
2	■	■	■	■	■	■
3	■	■	■	■	■	■
4	■	■	■	■	■	■
5	■	■	■	■	■	■
6	■	■	■	■	■	■


# Tasks on graphs

## □ Graph classification




# Tasks on graphs

## □ Node classification




# Tasks on graphs

## □ Edge classification




# Graph convolution


## □ Convolution on a neighborhood




**CNN:** Pixel convolution




**CNN:** Pixel convolution  
(as a graph)




**GNN:** Graph convolution



# Graph convolution




Generating features from nodes 1 **hop** away  
**A,C, and D**




Now embedding for node **B** has  
information from **A,C, and D.**

# Graph convolution




Generating features from nodes **1 hop** away  
**A,C, and D**




But we don't want to forget  
information about B either.

Now embedding for node B has  
information from A,C,D, **and B  
itself.**

# Two layers of graph convolution




Including features from nodes **2 hops** away  
**A,B,C,D,E, and F**




# Receptive fields in graph neural networks

- Increasing receptive fields with more layers/hops



# Graph convolution



# Graph convolution

- **Recap: Simple neighborhood aggregation:**

$$\mathbf{h}_v^k = \sigma \left( \mathbf{W}_k \sum_{u \in N(v)} \frac{\mathbf{h}_u^{k-1}}{|N(v)|} + \mathbf{B}_k \mathbf{h}_v^{k-1} \right)$$

- Graph convolutional operator:

- Aggregates messages across neighborhoods,  $N(v)$
- $\alpha_{vu} = 1/|N(v)|$  is the **weighting factor (importance)** of node  $u$ 's message to node  $v$
- $\Rightarrow \alpha_{vu}$  is defined **explicitly** based on the structural properties of the graph
- $\Rightarrow$  All neighbors  $u \in N(v)$  are equally important to node  $v$

# Graph attention network

Can we do better than simple neighborhood aggregation?

Can we let weighting factors  $\alpha_{vu}$  to be implicitly defined?

- **Goal:** Specify arbitrary importances to different neighbors of each node in the graph
- **Idea:** Compute embedding  $h$  of each node in the graph following an **attention strategy**:
  - Nodes attend over their neighborhoods' message
  - Implicitly specifying different weights to different nodes in a neighborhood

# Graph attention

- Let  $\alpha_{vu}$  be computed as a byproduct of an **attention mechanism  $a$** :
  - Let  $a$  compute **attention coefficients  $e_{vu}$**  across pairs of nodes  $u, v$  based on their messages:
$$e_{vu} = a(\mathbf{W}_k \mathbf{h}_u^{k-1}, \mathbf{W}_k \mathbf{h}_v^{k-1})$$
    - $e_{vu}$  indicates the importance of node  $u$ 's message to node  $v$
  - Normalize coefficients** using the softmax function in order to be comparable across different neighborhoods:

$$\alpha_{vu} = \frac{\exp(e_{vu})}{\sum_{k \in N(v)} \exp(e_{vk})}$$

$$\mathbf{h}_v^k = \sigma\left(\sum_{u \in N(v)} \alpha_{vu} \mathbf{W}_k \mathbf{h}_u^{k-1}\right)$$

**Next:** What is the form of attention mechanism  $a$ ?



An application of Graph Neural Network  
- SuperGlue feature matching



# SuperGlue: Learning Feature Matching with Graph Neural Networks



Paul-Edouard Sarlin<sup>1</sup>  
Tomasz Malisiewicz<sup>2</sup>


Daniel DeTone<sup>2</sup>  
Andrew Rabinovich<sup>2</sup>

**ETH** zürich




# The importance of context

no  
SuperGlue



with  
SuperGlue



# Problem formulation



Inputs



Outputs





- Images **A** and **B**
- **2 sets of  $M, N$  local features**
  - Keypoints:  $\mathbf{p}_i := (x, y, c)_i$ 
    - Coordinates  $(x, y)$
    - Confidence  $c$
  - Visual descriptors:  $\mathbf{d}_i$



Single a match per keypoint  
+ occlusion and noise  
→ a **soft partial assignment**:

$$\mathbf{P} \in [0, 1]^{M \times N}$$





## A Graph Neural Network with attention


Encodes **contextual cues** & priors

**Reasons** about the 3D scene

## Solving a partial assignment problem

Differentiable **solver**


Enforces the assignment constraints  
= **domain knowledge**



- Initial representation for each keypoints  $i$  :  ${}^{(0)}\mathbf{x}_i$
- Combines visual appearance and position with an MLP:

$${}^{(0)}\mathbf{x}_i = \mathbf{d}_i + \text{MLP}(\mathbf{p}_i)$$

Multi-Layer Perceptron




**Update** the representation based on other keypoints:

- in the same image: “**self**” edges
- in the other image: “**cross**” edges

$$(\ell) \mathbf{x}_i^A \longrightarrow (\ell+1) \mathbf{x}_i^A$$

→ A complete **graph** with two types of edges



**Update the representation using a Message Passing Neural Network**


$${}^{(\ell+1)}\mathbf{x}_i^A = {}^{(\ell)}\mathbf{x}_i^A + \text{MLP} \left( \left[ {}^{(\ell)}\mathbf{x}_i^A \parallel \mathbf{m}_{\mathcal{E} \rightarrow i} \right] \right)$$

↑  
the message


# Attentional Aggregation

- Compute the **message**  $\mathbf{m}_{\mathcal{E} \rightarrow i}$  using **self** and **cross attention**
- Soft database retrieval: query  $\mathbf{q}_i$ , key  $\mathbf{k}_j$ , and value  $\mathbf{v}_j$

$$\mathbf{m}_{\mathcal{E} \rightarrow i} = \sum_{j:(i,j) \in \mathcal{E}} \alpha_{ij} \mathbf{v}_j \quad \left| \quad \begin{aligned} \mathbf{q}_i &= \mathbf{W}_1 \mathbf{x}_i^{(\ell)} + \mathbf{b}_1 \\ \begin{bmatrix} \mathbf{k}_j \\ \mathbf{v}_j \end{bmatrix} &= \begin{bmatrix} \mathbf{W}_2 \\ \mathbf{W}_3 \end{bmatrix} \mathbf{x}_j^{(\ell)} + \begin{bmatrix} \mathbf{b}_2 \\ \mathbf{b}_3 \end{bmatrix} \end{aligned} \right.$$
$$\alpha_{ij} = \text{Softmax}_j (\mathbf{q}_i^\top \mathbf{k}_j)$$





**Self-attention**  
= intra-image  
information  
flow



**Cross-attention**  
= inter-image

Attention builds a  
**soft, dynamic,**  
**sparse graph**






Compute a **score matrix**  $S \in \mathbb{R}^{M \times N}$   
for all matches:


$$\mathbf{f}_i^A = \mathbf{W} \cdot {}^{(L)}\mathbf{x}_i^A + \mathbf{b}$$

$$S_{i,j} = \langle \mathbf{f}_i^A, \mathbf{f}_j^B \rangle$$






# Results: indoor - ScanNet

SuperPoint + NN + heuristics




SuperPoint + SuperGlue






SuperGlue: more **correct matches** and fewer **mismatches**

# Results: outdoor - SfM

SuperPoint + NN + OA-Net (inlier classifier)



SuperPoint + SuperGlue



SuperGlue: more **correct matches** and fewer **mismatches**